



Sidebar▼

[Home](#) → [Tutorials](#) → [Linux](#) → [Device Drivers](#) → **Linux Device Driver Tutorial Part 14 - Workqueue in Linux Kernel Part 1**

Device Drivers



Linux Device Driver Tutorial Part 14 - Workqueue in Linux Kernel Part 1

X

carries on the discussion on character drivers and their implementation. This is Part 14 of the Linux device driver tutorial. In our previous tutorial, we have seen the [Example of Interrupt](#) through Device Driver Programming. Now we will see one of the Bottomhalf which is Workqueue in Linux Kernel.

Apple Music Turns 5 as It Continues Rivalry With Spotify

Post Contents [\[hide\]](#)

- 1 Bottom Half
- 2 Workqueue in Linux Kernel
- 3 Using Global Workqueue (Global Worker Thread)
 - 3.1 Initialize work using Static Method
 - 3.1.1 Example
 - 3.2 Schedule work to the Workqueue
 - 3.2.1 Schedule_work
 - 3.2.2 Scheduled_delayed_work
 - 3.2.3 Schedule_work_on
 - 3.2.4 Scheduled_delayed_work_on
 - 3.3 Delete work from workqueue
 - 3.4 Cancel Work from workqueue
 - 3.5 Check the workqueue

X

[4.0.1 Share this:](#)[4.0.2 Like this:](#)[4.0.3 Related](#)

Bottom Half

When Interrupt triggers, Interrupt Handler should be executed very quickly and it should not run for more time (it should not perform time-consuming tasks). If we have the interrupt handler which is doing more tasks then we need to divide into two halves.

1. Top Half
2. Bottom Half

Top Half is nothing but our interrupt handler. If our interrupt handler is doing less task, then the top half is more than enough. No need for the bottom half in that situation. But if we have more work to do when interrupt hits, then we need bottom half. The bottom half runs in the

A circular icon with a black 'X' inside, typically used as a close or delete button in user interfaces.

There are 4 bottom half mechanisms are available in Linux:

1. [Workqueue](#)
2. Threaded IRQ
3. Softirq
4. [Tasklet](#)

In this tutorial, we will see Workqueue in Linux Kernel.

Workqueue in Linux Kernel

Work queues are added in the Linux kernel 2.6 version. Work queues are a different form of deferring work. Work queues defer work into a kernel thread; this bottom half always runs in process context. Because workqueue is allowing users to create a kernel thread and bind work to the kernel thread. So, this will run in process context and the work queue can sleep.

- Code deferred to a work queue has all the usual benefits of process context.
- Most importantly, work queues are schedulable and can therefore sleep.

Normally, it is easy to decide between using workqueue and [softirq/tasklet](#).

X

There are two ways to implement Workqueue in Linux kernel.

1. Using global workqueue
2. Creating Own workqueue (We will see in the [next tutorial](#))

Using Global Workqueue (Global Worker Thread)

In this tutorial, we will focus on this method.

In this method no need to create any workqueue or worker thread. So in this method, we only need to initialize work. We can initialize the work using two methods.

- Static method
- Dynamic method (We will see in the [next tutorial](#))

Initialize work using Static Method

The below call creates a workqueue by the name and the function that we are passing in the second argument gets scheduled in the queue.

```
DECLARE_WORK(name, void (*func)(void *))
```

Where,

name: The name of the “work_struct” structure that has to be created.

func: The function to be scheduled in this workqueue.

Example

```
1 DECLARE_WORK(workqueue,workqueue_fn);
```

Schedule work to the Workqueue

These below functions used to allocate the work to the queue.

```
int schedule_work( struct work_struct *work );
```

where,

work – job to be done

Returns zero if *work* was already on the kernel-global workqueue and non-zero otherwise.

Scheduled_delayed_work

After waiting for a given time this function puts a job in the kernel-global workqueue.

```
int scheduled_delayed_work( struct delayed_work *dwork,  
unsigned long delay );
```

where,

dwork – job to be done

delay – number of jiffies to wait or 0 for immediate execution

Schedule_work_on

This puts a job on a specific CPU.

```
int schedule_work_on( int cpu, struct work_struct *work );
```

where,

cpu– CPU to put the work task on

work– job to be done

```
int scheduled_delayed_work_on(
    int cpu, struct delayed_work *dwork, unsigned long
    delay );
```

where,

cpu – CPU to put the work task on

dwork – job to be done

delay – number of jiffies to wait or 0 for immediate execution

Delete work from workqueue

There are also a number of helper functions that you can use to flush or cancel work on work queues. To flush a particular work item and block until the work is complete, you can make a call to **flush_work**. All work on a given work queue can be completed using a call to **flush_workqueue**. In both cases, the caller blocks until the operation are complete. To flush the kernel-global work queue, call **flush_scheduled_work**.

```
int flush_work( struct work_struct *work );
void flush_scheduled_work( void );
```

Cancel Work from workqueue

You can cancel work if it is not already executing in a handler. A call to **cancel_work_sync** will terminate the work in the queue or block until the callback has finished (if the work is already in progress in the handler). If the work is delayed, you can use a call to **cancel_delayed_work_sync**.

```
int cancel_work_sync( struct work_struct *work );
int cancel_delayed_work_sync( struct delayed_work *dwork );
```

Check the workqueue

```
work_pending( work );
delayed_work_pending( work );
```

Programming

Driver Source Code

I took the source code from the previous [interrupt example tutorial](#). In that source code, When we read the `/dev/etx_device` the interrupt will hit (To understand interrupts in Linux go to [this tutorial](#)). Whenever interrupt hits, I'm scheduling the work to the workqueue. I'm not going to do any job in both interrupt handler and workqueue function since it is a tutorial post. But in real workqueue, this function can be used to carry out any operations that need to be scheduled.

```
1  #include <linux/kernel.h>
2  #include <linux/init.h>
3  #include <linux/module.h>
4  #include <linux/kdev_t.h>
5  #include <linux/fs.h>
6  #include <linux/cdev.h>
7  #include <linux/device.h>
8  #include<linux/slab.h>           //kmalloc()
9  #include<linux/uaccess.h>       //copy_to/from_user()
10 #include<linux/sysfs.h>
11 #include<linux/kobject.h>
12 #include <linux/interrupt.h>
13 #include <asm/io.h>
14 #include <linux/workqueue.h>    // Required for workqueues
15
16
17 #define IRQ_NO 11
18
19
20 void workqueue_fn(struct work_struct *work);
21
22 /*Creating work by Static Method */
23 DECLARE_WORK(workqueue,workqueue_fn);
24
25 /*Workqueue Function*/
26 void workqueue_fn(struct work_struct *work)
27 {
28     printk(KERN_INFO "Executing Workqueue Function\n");
29 }
30
31
32 //Interrupt handler for IRQ 11.
```



```

40
41 volatile int etx_value = 0;
42
43
44 dev_t dev = 0;
45 static struct class *dev_class;
46 static struct cdev etx_cdev;
47 struct kobject *kobj_ref;
48
49 static int __init etx_driver_init(void);
50 static void __exit etx_driver_exit(void);
51
52 /***** Driver Fuctions *****/
53 static int etx_open(struct inode *inode, struct file *file);
54 static int etx_release(struct inode *inode, struct file *file);
55 static ssize_t etx_read(struct file *filp,
56                        char __user *buf, size_t len, loff_t * off);
57 static ssize_t etx_write(struct file *filp,
58                        const char *buf, size_t len, loff_t * off);
59
60 /***** Sysfs Fuctions *****/
61 static ssize_t sysfs_show(struct kobject *kobj,
62                          struct kobj_attribute *attr, char *buf);
63 static ssize_t sysfs_store(struct kobject *kobj,
64                          struct kobj_attribute *attr, const char *buf, size_t count);
65
66 struct kobj_attribute etx_attr = __ATTR(etx_value, 0660, sysfs_show, sysfs_store);
67
68 static struct file_operations fops =
69 {
70     .owner          = THIS_MODULE,
71     .read           = etx_read,
72     .write          = etx_write,
73     .open           = etx_open,
74     .release        = etx_release,
75 };
76
77 static ssize_t sysfs_show(struct kobject *kobj,
78                          struct kobj_attribute *attr, char *buf)
79 {
80     printk(KERN_INFO "Sysfs - Read!!!\n");
81     return sprintf(buf, "%d", etx_value);
82 }
83
84 static ssize_t sysfs_store(struct kobject *kobj,
85                          struct kobj_attribute *attr, const char *buf, size_t count)
86 {
87     printk(KERN_INFO "Sysfs - Write!!!\n");
88     sscanf(buf, "%d", &etx_value);
89     return count;
90 }
91
92 static int etx_open(struct inode *inode, struct file *file)
93 {
94     printk(KERN_INFO "Device File Opened...!!!\n");
95     return 0;

```

```

103
104 static ssize_t etx_read(struct file *filp,
105                        char __user *buf, size_t len, loff_t *off)
106 {
107     printk(KERN_INFO "Read function\n");
108     asm("int $0x3B"); // Corresponding to irq 11
109     return 0;
110 }
111 static ssize_t etx_write(struct file *filp,
112                        const char __user *buf, size_t len, loff_t *off)
113 {
114     printk(KERN_INFO "Write Function\n");
115     return 0;
116 }
117
118
119 static int __init etx_driver_init(void)
120 {
121     /*Allocating Major number*/
122     if((alloc_chrdev_region(&dev, 0, 1, "etx_Dev")) < 0){
123         printk(KERN_INFO "Cannot allocate major number\n");
124         return -1;
125     }
126     printk(KERN_INFO "Major = %d Minor = %d \n", MAJOR(dev), MINOR(dev));
127
128     /*Creating cdev structure*/
129     cdev_init(&etx_cdev, &fops);
130
131     /*Adding character device to the system*/
132     if((cdev_add(&etx_cdev, dev, 1)) < 0){
133         printk(KERN_INFO "Cannot add the device to the system\n");
134         goto r_class;
135     }
136
137     /*Creating struct class*/
138     if((dev_class = class_create(THIS_MODULE, "etx_class")) == NULL){
139         printk(KERN_INFO "Cannot create the struct class\n");
140         goto r_class;
141     }
142
143     /*Creating device*/
144     if((device_create(dev_class, NULL, dev, NULL, "etx_device")) == NULL){
145         printk(KERN_INFO "Cannot create the Device 1\n");
146         goto r_device;
147     }
148
149     /*Creating a directory in /sys/kernel/ */
150     kobj_ref = kobject_create_and_add("etx_sysfs", kernel_kobj);
151
152     /*Creating sysfs file for etx_value*/
153     if(sysfs_create_file(kobj_ref, &etx_attr.attr)){
154         printk(KERN_INFO "Cannot create sysfs file.....\n");
155         goto r_sysfs;
156     }
157     if (request_irq(IRQ_NO, irq_handler, IRQF_SHARED, "etx_device", (void *)0))
158         printk(KERN_INFO "my_device: cannot register IRQ ");

```

```

166
167 r_sysfs:
168     kobject_put(kobj_ref);
169     sysfs_remove_file(kernel_kobj, &etx_attr.attr);
170
171 r_device:
172     class_destroy(dev_class);
173 r_class:
174     unregister_chrdev_region(dev,1);
175     cdev_del(&etx_cdev);
176     return -1;
177 }
178
179 void __exit etx_driver_exit(void)
180 {
181     free_irq(IRQ_NO, (void *) (irq_handler));
182     kobject_put(kobj_ref);
183     sysfs_remove_file(kernel_kobj, &etx_attr.attr);
184     device_destroy(dev_class, dev);
185     class_destroy(dev_class);
186     cdev_del(&etx_cdev);
187     unregister_chrdev_region(dev, 1);
188     printk(KERN_INFO "Device Driver Remove...Done!!!\n");
189 }
190
191 module_init(etx_driver_init);
192 module_exit(etx_driver_exit);
193
194 MODULE_LICENSE("GPL");
195 MODULE_AUTHOR("EmbeTronicX <embetronicx@gmail.com or admin@embetronicx.com>");
196 MODULE_DESCRIPTION("A simple device driver - Workqueue part 1");
197 MODULE_VERSION("1.10");

```

MakeFile

```

1 obj-m += driver.o
2
3 KDIR = /lib/modules/$(shell uname -r)/build
4
5
6 all:
7     make -C $(KDIR) M=$(shell pwd) modules
8
9 clean:
10    make -C $(KDIR) M=$(shell pwd) clean

```

Building and Testing Driver

- Build the driver by using Makefile (***sudo make***)
- Load the driver using ***sudo insmod driver.ko***
- To trigger the interrupt read device file (***sudo cat /dev/etx_device***)

```
[11213.945181] Device Driver Insert...Done!!!  
[11217.255727] Device File Opened...!!!  
[11217.255747] Read function  
[11217.255783] Shared IRQ: Interrupt Occurred  
[11217.255845] Executing Workqueue Function  
[11217.255860] Device File Closed...!!!
```

- We can able to see the print **“Shared IRQ: Interrupt Occurred”** and **“Executing Workqueue Function”**
- Unload the module using **sudo rmmod driver**

In our [next tutorial](#), we will discuss Workqueue using the Dynamic method.


5

Article Rating

**Share this:**

Post

Tweet






Share

1

**Like this:**

Loading...

Related**Interrupts in
Linux**Linux Device Driver
Tutorial Part 12 -Tasklet Tutorial (Static Method)
Linux Device Driver Part 20Linux Device Driver
Tutorial Part 20 – Tasklet |**Interrupts in
Linux**Example Driver
ProgramLinux Device Driver
Tutorial Part 13 -

X

☒ Subscribe ▼

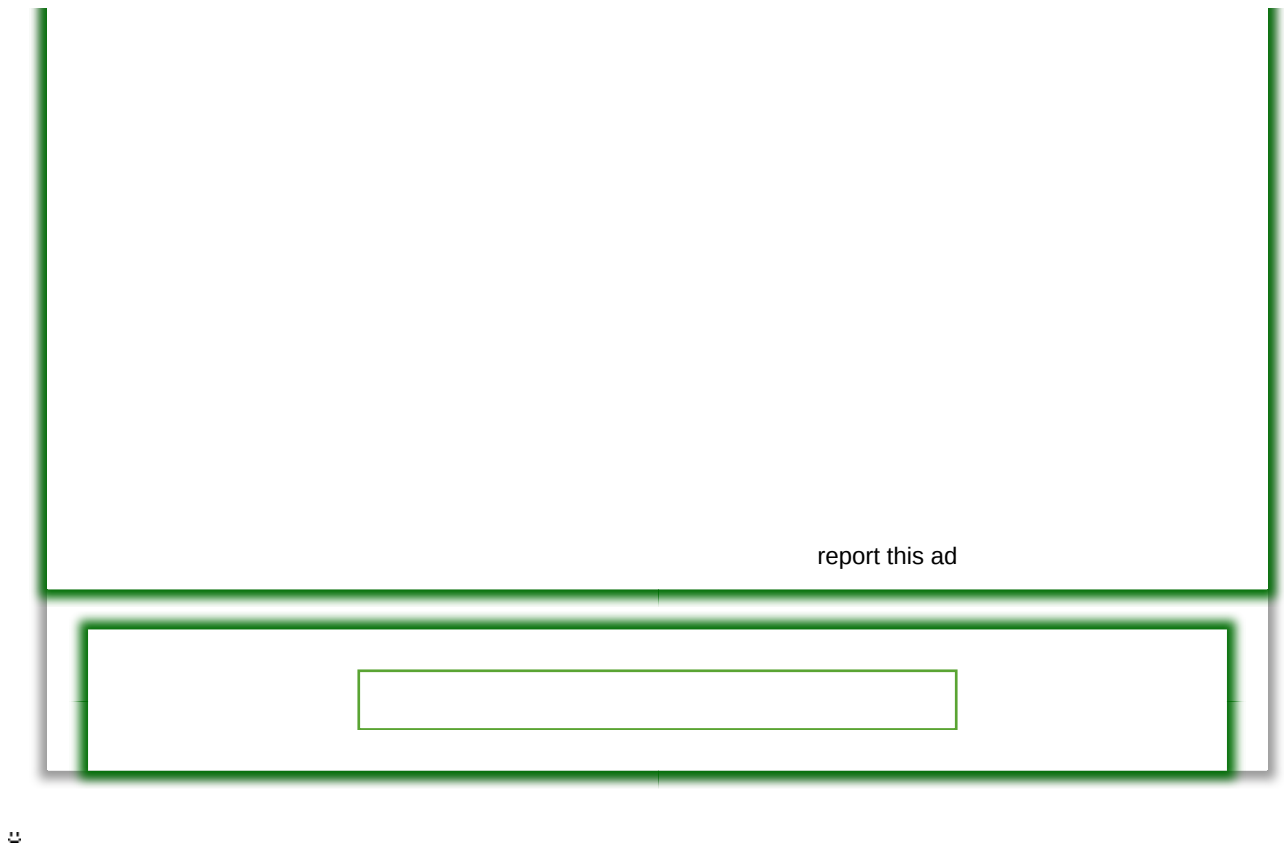
Connect with

[Login](#)*Be the First to Comment!***B** *I* U 

This site uses Akismet to reduce spam. [Learn how your comment data is processed.](#)

0 COMMENTS

X



u

