

[Sidebar](#) ▼

[Home](#) → [Tutorials](#) → [Linux](#) → [Device Drivers](#) → **Linux Device Driver Tutorial Part 16 - Workqueue in Linux Kernel Part 3**

📁 Device Drivers



Linux Device Driver Tutorial Part 16 - Workqueue in Linux Kernel Part 3

This ³ is the [Series on Linux Device Driver](#). The aim of this series is to provide easy and practical examples that anyone can understand. In our previous [tutorials](#), we have used global workqueue. But in this tutorial, we are going to use our own workqueue in the Linux device driver.

Apple Music Turns 5 as It Continues Rivalry With Spotify

Post Contents [\[hide\]](#)

- 1 [Workqueue in Linux Device Driver](#)
- 2 [Create and destroy workqueue structure](#)
 - 2.1 [alloc_workqueue](#)
 - 2.1.1 [WQ_* flags](#)
- 3 [Queuing Work to workqueue](#)
 - 3.1 [queue_work](#)
 - 3.2 [queue_work_on](#)
 - 3.3 [queue_delayed_work](#)
 - 3.4 [queue_delayed_work_on](#)
- 4 [Programming](#)
 - 4.1 [Driver Source Code](#)
 - 4.2 [MakeFile](#)
- 5 [Building and Testing Driver](#)
- 6 [Difference between Schedule_work and queue_work](#)
 - 6.0.1 [Share this:](#)
 - 6.0.2 [Like this:](#)
 - 3 6.0.3 [Related](#)

Workqueue in Linux Device Driver

In our previous ([Part 1](#), [Part 2](#)) tutorials we haven't created any of the workqueue. We were just creating work and scheduling that work to the global workqueue. Now we are going to create our own workqueue. Let's

get into the tutorial.

The core workqueue is represented by structure `struct workqueue_struct`, which is the structure onto which work is placed. This work is added to the queue in the top half (Interrupt context) and the execution of this work happened in the bottom half (Kernel context).

The work is represented by structure `struct work_struct`, which identifies the work and the deferral function.

Create and destroy workqueue structure

Workqueues are created through a macro called `create_workqueue`, which returns a `workqueue_struct` reference. You can remote this workqueue later (if needed) through a call to the `destroy_workqueue` function.

```
struct workqueue_struct *create_workqueue( name );
```

```
void destroy_workqueue( struct workqueue_struct * );
```

You should use `create_singlethread_workqueue()` for creating workqueue when you want to create only a single thread for all the processors.

Since `create_workqueue` and `create_singlethread_workqueue()` are macros. Both are using the `alloc_workqueue` function in the background.

```
1 #define create_workqueue(name)
2     alloc_workqueue("%S", WQ_MEM_RECLAIM, 1, (name))
3 #define create_singlethread_workqueue(name)
4     alloc_workqueue("%S", WQ_UNBOUND | WQ_MEM_RECLAIM, 1, (name))
```

alloc_workqueue

Allocate a workqueue with the specified parameters.

```
alloc_workqueue ( fmt, flags, max_active );
```

fmt - printf format for the name of the workqueue

flags - WQ_* flags

max_active – max in-flight work items, 0 for default

This will return Pointer to the allocated workqueue on success, **NULL** on failure.

WQ_* flags

This is the second argument of **alloc_workqueue**.

WQ_UNBOUND

Work items queued to an unbound **wq** are served by the special worker-pools which host workers who are not bound to any specific CPU. This makes the **wq** behave like a simple execution context provider without concurrency management. The unbound worker-pools try to start the execution of work items as soon as possible. Unbound **wq** sacrifices locality but is useful for the following cases.

- Wide fluctuation in the concurrency level requirement is expected and using bound **wq** may end up creating a large number of mostly unused workers across different CPUs as the issuer hops through different CPUs.
- Long-running CPU intensive workloads which can be better managed by the system scheduler.

WQ_FREEZABLE

A freezable **wq** participates in the freeze phase of the system suspend operations. Work items on the **wq** are drained and no new work item starts execution until thawed.

WQ_MEM_RECLAIM

All **wq** which might be used in the memory reclaim paths **MUST** have this flag set. The **wq** is guaranteed to have at least one execution context regardless of memory pressure.

WQ_HIGHPRI

Work items of a highpri wq are queued to the highpri worker-pool of the target CPU. Highpri worker-pools are served by worker threads with elevated nice levels.

Note that normal and highpri worker-pools don't interact with each other. Each maintains its separate pool of workers and implements concurrency management among its workers.

WQ_CPU_INTENSIVE

Work items of a CPU intensive wq do not contribute to the concurrency level. In other words, runnable CPU intensive work items will not prevent other work items in the same worker-pool from starting execution. This is useful for bound work items that are expected to hog CPU cycles so that their execution is regulated by the system scheduler.

Although CPU intensive work items don't contribute to the concurrency level, the start of their executions is still regulated by the concurrency management and runnable non-CPU-intensive work items can delay execution of CPU intensive work items.

This flag is meaningless for unbound wq.

Queuing Work to workqueue

With the work structure initialized, the next step is enqueueing the work on a workqueue. You can do this in a few ways.

queue_work

This will queue the work to the CPU on which it was submitted, but if the CPU dies it can be processed by another CPU.

```
3 int queue_work( struct workqueue_struct *wq, struct
    work_struct *work );
```

Where,

wq - workqueue to use

work – work to queue

It returns **false** if **work** was already on a queue, **true** otherwise.

queue_work_on

This puts work on a specific CPU.

```
int queue_work_on( int cpu, struct workqueue_struct *wq,
                  struct work_struct *work );
```

Where,

cpu– cpu to put the work task on

wq – workqueue to use

work– job to be done

queue_delayed_work

After waiting for a given time this function puts work in the workqueue.

```
int queue_delayed_work( struct workqueue_struct *wq,
                       struct delayed_work *dwork, unsigned long delay );
```

Where,

wq – workqueue to use

dwork – work to queue

delay – number of jiffies to wait before queueing or 0 for immediate execution

3

queue_delayed_work_on

After waiting for a given time this puts a job in the workqueue on the specified CPU.

```
int queue_delayed_work_on( int cpu, struct workqueue_struct
```

```

        *wq,
    struct delayed_work *dwork, unsigned long delay );

```

Where,

cpu - CPU to put the work task on

wq - workqueue to use

dwork - work to queue

delay - number of jiffies to wait before queueing or 0 for immediate execution

Programming Driver Source Code

In that source code, When we read the **/dev/etx_device**, interrupt will hit (To understand interrupts in Linux go to [this tutorial](#)). Whenever interrupt hits, I'm scheduling the work to the workqueue. I'm not going to do any job in both interrupt handler and workqueue function since it is a tutorial post. But in real workqueue, this function can be used to carry out any operations that need to be scheduled.

We have created workqueue "own_wq" in init function.

Let's go through the code.

```

1  #include <linux/kernel.h>
2  #include <linux/init.h>
3  #include <linux/module.h>
4  #include <linux/kdev_t.h>
5  #include <linux/fs.h>
6  #include <linux/cdev.h>
7  #include <linux/device.h>
8  #include <linux/slab.h>           //kmalloc()
9  #include <linux/uaccess.h>       //copy_to/from_user()
10 #include <linux/sysfs.h>
11 #include <linux/kobject.h>
12 #include <linux/interrupt.h>
13 #include <asm/io.h>
14 #include <linux/workqueue.h>     // Required for workqueues
15
16
17 #define IRQ_NO 11
18
19 static struct workqueue_struct *own_workqueue;
20

```

```

21 static void workqueue_fn(struct work_struct *work);
22
23 static DECLARE_WORK(work, workqueue_fn);
24
25
26 /*Workqueue Function*/
27 static void workqueue_fn(struct work_struct *work)
28 {
29     printk(KERN_INFO "Executing Workqueue Function\n");
30     return;
31 }
32
33
34
35 //Interrupt handler for IRQ 11.
36 static irqreturn_t irq_handler(int irq,void *dev_id) {
37     printk(KERN_INFO "Shared IRQ: Interrupt Occurred\n");
38     /*Allocating work to queue*/
39     queue_work(own_workqueue, &work);
40
41     return IRQ_HANDLED;
42 }
43
44
45 volatile int etx_value = 0;
46
47
48 dev_t dev = 0;
49 static struct class *dev_class;
50 static struct cdev etx_cdev;
51 struct kobject *kobj_ref;
52
53 static int __init etx_driver_init(void);
54 static void __exit etx_driver_exit(void);
55
56 /***** Driver Fuctions *****/
57 static int etx_open(struct inode *inode, struct file *file);
58 static int etx_release(struct inode *inode, struct file *file);
59 static ssize_t etx_read(struct file *filp,
60     char __user *buf, size_t len,loff_t * off);
61 static ssize_t etx_write(struct file *filp,
62     const char *buf, size_t len, loff_t * off);
63
64 /***** Sysfs Fuctions *****/
65 static ssize_t sysfs_show(struct kobject *kobj,
66     struct kobj_attribute *attr, char *buf);
67 static ssize_t sysfs_store(struct kobject *kobj,
68     struct kobj_attribute *attr,const char *buf, size_t count);
69
70 struct kobj_attribute etx_attr = __ATTR(etx_value, 0660, sysfs_show, sysfs_store);
71
72 static struct file_operations fops =
73 {
74     .owner          = THIS_MODULE,
75     .read           = etx_read,
76     .write          = etx_write,
77     .open           = etx_open,
78     .release        = etx_release,
79 };
80
81 static ssize_t sysfs_show(struct kobject *kobj,
82     struct kobj_attribute *attr, char *buf)
83 {

```



```

84     printk(KERN_INFO "Sysfs - Read!!!\n");
85     return sprintf(buf, "%d", etx_value);
86 }
87
88 static ssize_t sysfs_store(struct kobject *kobj,
89                          struct kobj_attribute *attr, const char *buf, size_t count)
90 {
91     printk(KERN_INFO "Sysfs - Write!!!\n");
92     sscanf(buf, "%d", &etx_value);
93     return count;
94 }
95
96 static int etx_open(struct inode *inode, struct file *file)
97 {
98     printk(KERN_INFO "Device File Opened...!!!\n");
99     return 0;
100 }
101
102 static int etx_release(struct inode *inode, struct file *file)
103 {
104     printk(KERN_INFO "Device File Closed...!!!\n");
105     return 0;
106 }
107
108 static ssize_t etx_read(struct file *filp,
109                       char __user *buf, size_t len, loff_t *off)
110 {
111     printk(KERN_INFO "Read function\n");
112     asm("int $0x3B"); // Corresponding to irq 11
113     return 0;
114 }
115 static ssize_t etx_write(struct file *filp,
116                        const char __user *buf, size_t len, loff_t *off)
117 {
118     printk(KERN_INFO "Write Function\n");
119     return 0;
120 }
121
122
123 static int __init etx_driver_init(void)
124 {
125     /*Allocating Major number*/
126     if((alloc_chrdev_region(&dev, 0, 1, "etx_Dev")) < 0){
127         printk(KERN_INFO "Cannot allocate major number\n");
128         return -1;
129     }
130     printk(KERN_INFO "Major = %d Minor = %d \n", MAJOR(dev), MINOR(dev));
131
132     /*Creating cdev structure*/
133     cdev_init(&etx_cdev, &fops);
134
135     /*Adding character device to the system*/
136     if((cdev_add(&etx_cdev, dev, 1)) < 0){
137         printk(KERN_INFO "Cannot add the device to the system\n");
138         goto r_class;
139     }
140
141     /*Creating struct class*/
142     if((dev_class = class_create(THIS_MODULE, "etx_class")) == NULL){
143         printk(KERN_INFO "Cannot create the struct class\n");
144         goto r_class;
145     }
146

```

```

147     /*Creating device*/
148     if((device_create(dev_class,NULL,dev,NULL,"etx_device")) == NULL){
149         printk(KERN_INFO "Cannot create the Device 1\n");
150         goto r_device;
151     }
152
153     /*Creating a directory in /sys/kernel/ */
154     kobj_ref = kobject_create_and_add("etx_sysfs",kernel_kobj);
155
156     /*Creating sysfs file for etx_value*/
157     if(sysfs_create_file(kobj_ref,&etx_attr.attr)){
158         printk(KERN_INFO "Cannot create sysfs file.....\n");
159         goto r_sysfs;
160     }
161     if (request_irq(IRQ_NO, irq_handler, IRQF_SHARED, "etx_device", (void *)(&irq_handler)))
162         printk(KERN_INFO "my_device: cannot register IRQ \n");
163         goto irq;
164     }
165
166     /*Creating workqueue */
167     own_workqueue = create_workqueue("own_wq");
168
169     printk(KERN_INFO "Device Driver Insert...Done!!!\n");
170     return 0;
171
172 irq:
173     free_irq(IRQ_NO,(void *)(&irq_handler));
174
175 r_sysfs:
176     kobject_put(kobj_ref);
177     sysfs_remove_file(kernel_kobj, &etx_attr.attr);
178
179 r_device:
180     class_destroy(dev_class);
181 r_class:
182     unregister_chrdev_region(dev,1);
183     cdev_del(&etx_cdev);
184     return -1;
185 }
186
187 void __exit etx_driver_exit(void)
188 {
189     /* Delete workqueue */
190     destroy_workqueue(own_workqueue);
191     free_irq(IRQ_NO,(void *)(&irq_handler));
192     kobject_put(kobj_ref);
193     sysfs_remove_file(kernel_kobj, &etx_attr.attr);
194     device_destroy(dev_class,dev);
195     class_destroy(dev_class);
196     cdev_del(&etx_cdev);
197     unregister_chrdev_region(dev, 1);
198     printk(KERN_INFO "Device Driver Remove...Done!!!\n");
199 }
200
201 module_init(etx_driver_init);
202 module_exit(etx_driver_exit);
203
204 MODULE_LICENSE("GPL");
205 MODULE_AUTHOR("EmbeTronicX <embetronicx@gmail.com>");
206 MODULE_DESCRIPTION("A simple device driver - Workqueue part 3");
207 MODULE_VERSION("1.12");

```

MakeFile

```

1 obj-m += driver.o
2
3 KDIR = /lib/modules/$(shell uname -r)/build
4
5
6 all:
7     make -C $(KDIR) M=$(shell pwd) modules
8
9 clean:
10    make -C $(KDIR) M=$(shell pwd) clean

```

Building and Testing Driver

- Build the driver by using Makefile (**sudo make**)
- Load the driver using **sudo insmod driver.ko**
- To trigger the interrupt read device file (**sudo cat /dev/etx_device**)
- Now see the Dmesg (**dmesg**)

```

[ 2562.609446] Major = 246 Minor = 0
[ 2562.649362] Device Driver Insert...Done!!!
[ 2565.133204] Device File Opened...!!!
[ 2565.133225] Read function
[ 2565.133248] Shared IRQ: Interrupt Occurred
[ 2565.133267] Executing Workqueue Function
[ 2565.140284] Device File Closed...!!!

```

- We can able to see the print “**Shared IRQ: Interrupt Occurred**” and “**Executing Workqueue Function**”
- Use “**ps -aef**” command to see our workqueue. You can able to see our workqueue which is “**own_wq**”

| UID | PID | PPID | C | STIME | TTY | TIME | CMD |
|------|------|------|---|-------|-------|------|-------------------|
| root | 3516 | 2 | | 0 | 21:35 | ? | 00:00:00 [own_wq] |

- Unload the module using **sudo rmmod driver**

Difference between Schedule_work and queue_work

- If you want to use your own dedicated workqueue you should create workqueue using **create_workqueue**. In that time you need to put work on your workqueue by using **queue_work** function.
- If you don't want to create any own workqueue, you can use kernel global workqueue. In that condition, you can use **schedule_work** function to put your work to global workqueue.

In our [next tutorial](#), we will discuss linked list in the Linux device driver.



Article Rating




Share this:




Post

Tweet



Share


2



Like this:

Loading...


Related



Workqueue in
Linux Device Driver
Static Method




Linux Device Driver
Tutorial Part 14 -
Workqueue in Linux
Kernel Part 1
In "Device Drivers"



Workqueue in
Linux Device Driver
Dynamic Method



Linux Device Driver
Tutorial Part 15 -
Workqueue in Linux
Kernel Part 2
In "Device Drivers"



Tasklet Tutorial (Static Method)
Linux Device Driver Part 20



Linux Device Driver
Tutorial Part 20 - Tasklet |
Static Method
In "Device Drivers"


Connect with | [Login](#)



3 COMMENTS



Oldest ▼



swapna

February 19, 2018 3:33 AM

Good job, easy to understand the tutorial. Please keep posting for the following topics

softirq

tasklet

timer


Loading...



0



Reply



arun

September 18, 2018 1:34 AM

this error is occurred plz help on this

initialization from incompatible pointer type

[-Werror=incompatible-pointer-types]

.func = (f),

^

./include/linux/workqueue.h:184:25: note: in expansion of macro

'__WORK_INITIALIZER'

struct work_struct n = __WORK_INITIALIZER(n, f)

^

Loading...



0




Reply

3



EmbeTronicx India

 Reply to [arun](#)

September 19, 2018 1:04 AM

Hi Arun,
Please take the updated code arun.
Thanks.

Loading...

 0   Reply

[report this ad](#)

3

3