



Sidebar▼

[Home](#) → [Tutorials](#) → [Linux](#) → [Device Drivers](#) → **Linux Device Driver Tutorial Part 23 - Spinlock in Linux Kernel Part 1**

## Device Drivers



# **Linux Device Driver Tutorial Part 23 - Spinlock in Linux Kernel Part 1**

2

X

provide easy and practical examples that anyone can understand. This is the Linux Device Driver Tutorial Part 23 – Spinlock in Linux Kernel Part 1.

Apple Music Turns 5 as It Continues Rivalry With Spotify

### Post Contents [\[hide\]](#)

- 1 Prerequisites
- 2 Introduction
- 3 SpinLock
- 4 SpinLock in Linux Kernel Device Driver
  - 4.1 Initialize
    - 4.1.1 Static Method
    - 4.1.2 Dynamic Method
  - 4.2 Approach 1 (Locking between User context)
    - 4.2.1 Example
  - 4.3 Approach 2 (Locking between Bottom Halves)
  - 4.4 Approach 3 (Locking between User context and Bottom Halves)
    - 4.4.1 Example
  - 4.5 Approach 4 (Locking between Hard IRQ and Bottom Halves)
    - 2** 4.5.1 Example
  - 4.6 Approach 5 (Alternative way of Approach 4)
  - 4.7 Approach 6 (Locking between Hard IRQs)

X

[5.2.1 Share this:](#)[5.2.2 Like this:](#)[5.2.3 Related](#)

## Prerequisites

In the example section, I had used Kthread to explain Mutex. If you don't know what is Kthread and How to use it, then I would recommend you to explore that by using below link.

1. [Kthread Tutorial in Linux Kernel](#)
2. [Mutex Tutorial in Linux Kernel](#)

## Introduction

In our [previous tutorial](#), we have understood the use of Mutex and its Implementation. If you have understood Mutex then Spinlock is also similar. Both are used to protect a shared resource from being modified

A circular icon with a black 'X' inside, typically used as a close or exit button in user interfaces.

In the Mutex concept, when the thread is trying to lock or acquire the Mutex which is not available then that thread will go to sleep until that Mutex is available. Whereas in Spinlock it is different. The spinlock is a very simple single-holder lock. If a process attempts to acquire a spinlock and it is unavailable, the process will keep trying (spinning) until it can acquire the lock. This simplicity creates a small and fast lock.

Like Mutex, there are two possible states in Spinlock: **Locked** or **Unlocked**.

## SpinLock in Linux Kernel Device Driver

If the kernel is running on a uniprocessor and **CONFIG\_SMP**, **CONFIG\_PREEMPT** aren't enabled while compiling the kernel then spinlock will not be available. Because there is no reason to have a lock when no one else can run at the same time.

But if you have disabled **CONFIG\_SMP** and enabled **CONFIG\_PREEMPT** then spinlock will simply disable preemption, which is sufficient to prevent any races.

## Initialize

We can initialize Spinlock in Linux kernel in two ways.

## Static Method

You can statically initialize a Spinlock using the macro given below.

```
DEFINE_SPINLOCK(etx_spinlock);
```

The macro given above will create a **spinlock\_t** variable in the name of **etx\_spinlock** and initialize to **UNLOCKED STATE**. Take a look at the expansion of **DEFINE\_SPINLOCK** below.

```
#define DEFINE_SPINLOCK(x)    spinlock_t x = __SPIN_LOCK_UNLOCKED(x)
```

## Dynamic Method

If you want to initialize dynamically you can use the method as given below.

```
1 spinlock_t etx_spinlock;  
2 spin_lock_init(&etx_spinlock);
```

You can use any one of the methods.

After initializing the spinlock, there are several ways to use spinlock to lock or unlock, based on where the spinlock is used; either in user context or interrupt context. Let's look at the approaches with these situations.

## Approach 1 (Locking between User context)

If you share data with user context (between Kernel Threads), then you can use this approach.

**Lock:**

2

```
spin_lock(spinlock_t *lock)
```

### **spin\_trylock(spinlock\_t \*lock)**

Locks the spinlock if it is not already locked. If unable to obtain the lock it exits with an error and do not spin. It **returns** non-zero if it obtains the lock otherwise returns zero.

#### **Unlock:**

### **spin\_unlock(spinlock\_t \*lock)**

It does the reverse of the lock. It will unlock which is locked by the above call.

#### **Checking Lock:**

### **spin\_is\_locked(spinlock\_t \*lock)**

This is used to check whether the lock is available or not. It **returns** non-zero if the lock is currently acquired. otherwise returns zero.

## **Example**

```
1 //Thread 1
2 int thread_function1(void *pv)
3 {
4     while(!kthread_should_stop()) {
5         spin_lock(&etx_spinlock);
6         etx_global_variable++;
7         printk(KERN_INFO "In EmbeTronicX Thread Function1 %lu\n", etx_global_variable);
8         spin_unlock(&etx_spinlock);
9         msleep(1000);
10    }
11    return 0;
12 }
13
14 //Thread 2
15 int thread_function2(void *pv)
```

```
23     }  
24     return 0;  
25 }
```

## Approach 2 (Locking between Bottom Halves)

If you want to share data between two different Bottom halves or the same bottom halves, then you can use the [Approach 1](#).

## Approach 3 (Locking between User context and Bottom Halves)

If you share data with a [bottom half](#) and user context (like Kernel Thread), then this approach will be useful.

### Lock:

```
spin_lock_bh(spinlock_t *lock)
```

It disables soft interrupts on that CPU, then grabs the lock. This has the effect of preventing softirqs, tasklets, and bottom halves from running on the local CPU. Here the suffix '**\_bh**' refers to "**Bottom Halves**".

### Unlock:

```
spin_unlock_bh(spinlock_t *lock)
```

It will release the lock and re-enables the soft interrupts which are disabled by the above call.

## Example

```
1 //Thread  
2 int thread_function(void *pv)  
3 {  
4     while(!kthread_should_stop()) {  
5         spin_lock_bh(&etx_spinlock);  
6         etx_global_variable++;  
7     }
```

X

```
14 void tasklet_fn(unsigned long arg)
15 {
16     spin_lock_bh(&etx_spinlock);
17     etx_global_variable++;
18     printk(KERN_INFO "Executing Tasklet Function : %lu\n", etx_global_variable);
19     spin_unlock_bh(&etx_spinlock);
20 }
```

## Approach 4 (Locking between Hard IRQ and Bottom Halves)

If you share data between Hardware ISR and Bottom halves then you have to disable the IRQ before lock. Because the bottom halves processing can be interrupted by a hardware interrupt. So this will be used in that scenario.

### Lock:

```
spin_lock_irq(spinlock_t *lock)
```

This will disable interrupts on that cpu, then grab the lock.

### Unlock:

```
spin_unlock_irq(spinlock_t *lock)
```

2

It will release the lock and re-enables the interrupts which are disabled

X



```
1  /*Tasklet Function*/
2  void tasklet_fn(unsigned long arg)
3  {
4      spin_lock_irq(&etx_spinlock);
5      etx_global_variable++;
6      printk(KERN_INFO "Executing Tasklet Function : %lu\n", etx_global_variable);
7      spin_unlock_irq(&etx_spinlock);
8  }
9
10 //Interrupt handler for IRQ 11.
11 static irqreturn_t irq_handler(int irq,void *dev_id) {
12     spin_lock_irq(&etx_spinlock);
13     etx_global_variable++;
14     printk(KERN_INFO "Executing ISR Function : %lu\n", etx_global_variable);
15     spin_unlock_irq(&etx_spinlock);
16     /*Scheduling Task to Tasklet*/
17     tasklet_schedule(tasklet);
18     return IRQ_HANDLED;
19 }
```

## Approach 5 (Alternative way of Approach 4)

If you want to use a different variant rather than using `spin_lock_irq()` and `spin_unlock_irq()` then you can use this approach.

### Lock:

```
spin_lock_irqsave( spinlock_t *lock, unsigned long flags );
```

This will save whether interrupts were on or off in a `flags` word and grab the lock.

### Unlock:

```
spin_unlock_irqrestore( spinlock_t *lock, unsigned long flags );
```

This will release the spinlock and restores the interrupts using the `flags` argument.

2

## Approach 6 (Locking between Hard IRQs)

## programming

This code snippet explains how to create two threads that access a global variable (**etx\_gloabl\_variable**). So before accessing the variable, it should lock the spinlock. After that, it will release the spinlock. This example is using [Approach 1](#).

## Driver Source Code

```

1  #include <linux/kernel.h>
2  #include <linux/init.h>
3  #include <linux/module.h>
4  #include <linux/kdev_t.h>
5  #include <linux/fs.h>
6  #include <linux/cdev.h>
7  #include <linux/device.h>
8  #include <linux/slab.h>           //kmalloc()
9  #include <linux/uaccess.h>       //copy_to/from_user()
10 #include <linux/kthread.h>       //kernel threads
11 #include <linux/sched.h>         //task_struct
12 #include <linux/delay.h>
13
14 DEFINE_SPINLOCK(etx_spinlock);
15 //spinlock_t etx_spinlock;
16 unsigned long etx_global_variable = 0;
17
18 dev_t dev = 0;
19 static struct class *dev_class;
20 static struct cdev etx_cdev;
21
22 static int __init etx_driver_init(void);
23 static void __exit etx_driver_exit(void);
24
25 static struct task_struct *etx_thread1;
26 static struct task_struct *etx_thread2;
27
28 /***** Driver Fuctions *****/
29 static int etx_open(struct inode *inode, struct file *file);
30 static int etx_release(struct inode *inode, struct file *file);
31 static ssize_t etx_read(struct file *filp,
32                        char __user *buf, size_t len, loff_t * off);
33 static ssize_t etx_write(struct file *filp,
34                        const char *buf, size_t len, loff_t * off);
35 /*****/
36
37 int thread_function1(void *pv);
38 int thread_function2(void *pv);
39
40 int thread_function1(void *pv)
41 {
42

```

```
49         printk(KERN_INFO "Spinlock is locked in Thread Function1\n");
50     }
51     etx_global_variable++;
52     printk(KERN_INFO "In EmbeTronicX Thread Function1 %lu\n", etx_global_variab
53     spin_unlock(&etx_spinlock);
54     msleep(1000);
55 }
56 return 0;
57 }
58
59 int thread_function2(void *pv)
60 {
61     while(!kthread_should_stop()) {
62         spin_lock(&etx_spinlock);
63         etx_global_variable++;
64         printk(KERN_INFO "In EmbeTronicX Thread Function2 %lu\n", etx_global_variab
65         spin_unlock(&etx_spinlock);
66         msleep(1000);
67     }
68     return 0;
69 }
70
71 static struct file_operations fops =
72 {
73     .owner          = THIS_MODULE,
74     .read           = etx_read,
75     .write          = etx_write,
76     .open           = etx_open,
77     .release        = etx_release,
78 };
79
80 static int etx_open(struct inode *inode, struct file *file)
81 {
82     printk(KERN_INFO "Device File Opened...!!!\n");
83     return 0;
84 }
85
86 static int etx_release(struct inode *inode, struct file *file)
87 {
88     printk(KERN_INFO "Device File Closed...!!!\n");
89     return 0;
90 }
91
92 static ssize_t etx_read(struct file *filp,
93                        char __user *buf, size_t len, loff_t *off)
94 {
95     printk(KERN_INFO "Read function\n");
96     return 0;
97 }
98
99 static ssize_t etx_write(struct file *filp,
100                        const char __user *buf, size_t len, loff_t *off)
101 {
102     printk(KERN_INFO "Write Function\n");
103     return len;
104 }
```

```
112     }
113     printk(KERN_INFO "Major = %d Minor = %d \n", MAJOR(dev), MINOR(dev));
114
115     /*Creating cdev structure*/
116     cdev_init(&etx_cdev, &fops);
117
118     /*Adding character device to the system*/
119     if((cdev_add(&etx_cdev, dev, 1)) < 0){
120         printk(KERN_INFO "Cannot add the device to the system\n");
121         goto r_class;
122     }
123
124     /*Creating struct class*/
125     if((dev_class = class_create(THIS_MODULE, "etx_class")) == NULL){
126         printk(KERN_INFO "Cannot create the struct class\n");
127         goto r_class;
128     }
129
130     /*Creating device*/
131     if((device_create(dev_class, NULL, dev, NULL, "etx_device")) == NULL){
132         printk(KERN_INFO "Cannot create the Device \n");
133         goto r_device;
134     }
135
136
137     /* Creating Thread 1 */
138     etx_thread1 = kthread_run(thread_function1, NULL, "eTx Thread1");
139     if(etx_thread1) {
140         printk(KERN_ERR "Kthread1 Created Successfully...\n");
141     } else {
142         printk(KERN_ERR "Cannot create kthread1\n");
143         goto r_device;
144     }
145
146     /* Creating Thread 2 */
147     etx_thread2 = kthread_run(thread_function2, NULL, "eTx Thread2");
148     if(etx_thread2) {
149         printk(KERN_ERR "Kthread2 Created Successfully...\n");
150     } else {
151         printk(KERN_ERR "Cannot create kthread2\n");
152         goto r_device;
153     }
154     //spin_lock_init(&etx_spinlock);
155
156     printk(KERN_INFO "Device Driver Insert...Done!!!\n");
157     return 0;
158
159
160 r_device:
161     class_destroy(dev_class);
162 r_class:
163     unregister_chrdev_region(dev, 1);
164     cdev_del(&etx_cdev);
165     return -1;
166 }
167
```

```
175         unregister_chrdev_region(dev, 1);
176         printk(KERN_INFO "Device Driver Remove...Done!!\n");
177     }
178
179     module_init(etx_driver_init);
180     module_exit(etx_driver_exit);
181
182     MODULE_LICENSE("GPL");
183     MODULE_AUTHOR("EmbeTronicX <embetronicx@gmail.com>");
184     MODULE_DESCRIPTION("A simple device driver - Spinlock");
185     MODULE_VERSION("1.18");
```

## MakeFile

```
1 obj-m += driver.o
2
3 KDIR = /lib/modules/$(shell uname -r)/build
4
5 all:
6     make -C $(KDIR) M=$(shell pwd) modules
7
8 clean:
9     make -C $(KDIR) M=$(shell pwd) clean
```

In our [next part of the tutorial](#), we will see another variant of the spinlock (**Reader/writer spinlocks**).



### Share this:



### Like this:

Loading...



Linux Device Driver  
Tutorial Part 22 - Mutex in  
Linux Kernel  
In "Device Drivers"




Linux Device Driver  
Tutorial Part 24 - Read  
Write Spinlock in Linux  
Kernel (Spinlock Part 2)  
In "Device Drivers"



Linux Device Driver  
Tutorial Part 31 - Seqlock  
in Linux Kernel  
In "Device Drivers"

Connect with



**B** *I* U ~~S~~ <sup>1</sup>/<sub>2</sub>/<sub>3</sub> ≡ ≡≡ ” </> ∞ { } [ + ]



## 2 COMMENTS



Oldest ▼



May 23, 2019 6:31 AM

Loading...



0



↩ Reply



May 23, 2019 6:50 AM

2

During ISR execution interrupts on particular core are already disabled... So Spin lock is enough.....



