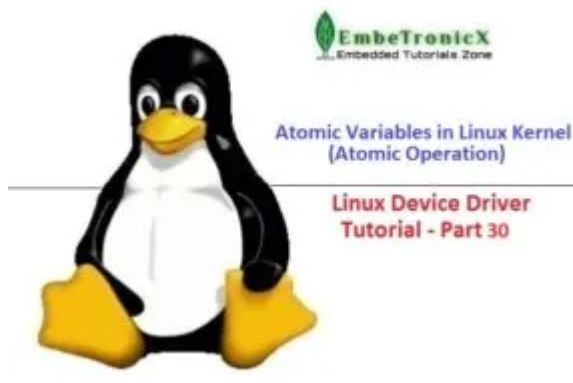




Sidebar▼

[Home](#) → [Tutorials](#) → [Linux](#) → [Device Drivers](#) → **Linux Device Driver Tutorial Part 30 - Atomic variable in Linux Device Driver**

📁 Device Drivers



Linux Device Driver Tutorial Part 30 - Atomic variable in Linux Device Driver

This is the [Series on Linux Device Driver](#). The aim of this series is to provide easy and practical examples that anyone can understand. This is the Linux Device Driver Tutorial Part 30 - atomic variable in Linux Device Driver (atomic operations).

1

Apple Music Turns 5 as It Continues Rivalry With Spotify

Post Contents [\[hide\]](#)

- 1 atomic variables in Linux Device Driver
- 2 Prerequisites
- 3 Introduction
- 4 atomic variables
- 5 Types of atomic variables
 - 5.1 Atomic Integer Operations
 - 5.1.1 Creating atomic variables
 - 5.1.2 Reading atomic variables
 - 5.1.2.1 atomic_read
 - 5.1.3 Other operations on atomic variables
 - 5.1.3.1 atomic_set
 - 5.1.3.2 atomic_add
 - 5.1.3.3 atomic_sub
 - 5.1.3.4 atomic_inc
 - 5.1.3.5 atomic_dec
 - 5.1.3.6 atomic_sub_and_test
 - 5.1.3.7 atomic_dec_and_test
 - 5.1.3.8 atomic_inc_and_test
 - 5.1.3.9 atomic_add_negative
 - 5.1.3.10 atomic_add_return
 - 5.1.3.11 atomic_add_unless
 - 5.2 Atomic Bitwise Operations

1

- 6 Example Programming
 - 6.1 Driver Source Code
 - 6.2 MakeFile
 - 6.3 Share this:
 - 6.4 Like this:
 - 6.5 Related

atomic variables in Linux Device Driver Prerequisites

In the below mentioned posts, we are using spinlock and mutex for synchronization. I would recommend you to explore that by using the below link.

- [Mutex Tutorial in Linux Device Driver](#)
- [Spinlock in Linux Device Driver - Part 1](#)
- [Spinlock in Linux Device Driver - Part 2](#)

Introduction

Before looking into atomic variables, we will see one example.

I have one integer or long variable named **etx_global_variable** which is shared between two threads. Two threads are just incrementing the variable like below.

Thread 1:

```
1 etx_global_variable++; //Accessing the variable
```

Thread 2:

```
1 etx_global_variable++; //Accessing the variable
```

1 Now we will see how it is incrementing internally in each instruction when both threads are running concurrently. Assume the initial value of **etx_global_variable** is 0.

Thread 1	Thread 2
----------	----------

get the value of etx_global_variable from memory (0)	get the value of etx_global_variable from memory (0)
increment etx_global_variable (0 -> 1)	—
—	increment etx_global_variable (0 -> 1)
write back the etx_global_variable value to memory (1)	—
—	write back the etx_global_variable value to memory (1)

Now the value of **etx_global_variable** is 1 after the two threads are processed. Are we expecting the same value which is 1? Nope. We are expecting the value of **etx_global_variable** should be 2. It is not running as expected because the global variable is sharing between two concurrent threads. So we need to implement synchronization because both the threads are accessing (writing/reading) the variable. We can implement synchronization like below using locks.

Thread 1:

```
1 lock(); //spinlock or mutex
2 etx_global_variable++; //Accessing the variable
3 unlock();
```

Thread 2:

```
1 lock(); //spinlock or mutex
2 etx_global_variable++; //Accessing the variable
3 unlock();
```

1
After this synchronization, we will see how it is incrementing internally when both threads are running concurrently. Assume the initial value of **etx_global_variable** is 0.

Thread 1	Thread 2
-----------------	-----------------

lock()	—
get the value of etx_global_variable from memory (0)	lock() (it will be stuck here because the lock is already taken by thread 1)
increment etx_global_variable (0 -> 1)	—
write back the etx_global_variable value to memory (1)	—
unlock()	—
—	get the value of etx_global_variable from memory (1)
—	increment etx_global_variable (1 -> 2)
—	write back the etx_global_variable value to memory (2)
—	unlock()

This will be the one possibility to run. Another possibility is mentioned below.

Thread 1	Thread 2
—	lock()
lock() (it will be stuck here because the lock is already taken by thread 2)	get the value of etx_global_variable from memory (0)
—	increment etx_global_variable (0 -> 1)
—	write back the etx_global_variable value to memory (1)
—	unlock()

get the value of etx_global_variable from memory (1)	—
increment etx_global_variable (1 -> 2)	—
write back the etx_global_variable value to memory (2)	—
unlock()	—

Great. That's all. Now we are getting 2 in the two methods mentioned above. But have anyone thought anytime that, why these things are required for a single variable? Why don't we have an alternate method for a single variable? Yes, obviously we have an alternate mechanism for integer and long variables. That is **the atomic operation**. If you use mutex/spinlock for just a single variable, it will add overhead. In this tutorial, we gonna see the atomic variable, atomic operation, and its usage.

atomic variables

The read, write and arithmetic operations on the atomic variables will be done in one instruction without interrupt.

So again we will take the same example mentioned above to explain the atomic variable operations. When we use the atomic method, that will work like below.

Thread 1	Thread 2
get, increment and store etx_global_variable (0 -> 1)	—
—	get, increment and store etx_global_variable (1 -> 2)

1
and another possibility will be,

Thread 1	Thread 2
----------	----------

—	get, increment and store etx_global_variable (0 -> 1)
get, increment and store etx_global_variable (1 -> 2)	—

So the extra locking mechanism is not required when we are using atomic variables since the operation is happening in one machine instruction.

An **atomic_t** holds an **int** value and **atomic64_t** holds the **long** value on all supported architectures.

In Linux Kernel Version 2.6, the atomic variable has defined below.

```
1 typedef struct {
2     volatile int counter;
3 } atomic_t;
4
5 #ifdef CONFIG_64BIT
6 typedef struct {
7     volatile long counter;
8 } atomic64_t;
9 #endif
```

Then later, they have removed **volatile** and defined as below.

```
1 typedef struct {
2     int counter;
3 } atomic_t;
4
5 #ifdef CONFIG_64BIT
6 typedef struct {
7     long counter;
8 } atomic64_t;
9 #endif
```

You can read [here](#) why they have removed volatile.

Types of atomic variables

Two different atomic variables are there.

- Atomic variables who operates on Integers
- Atomic variables who operates on Individual Bits

Atomic Integer Operations

When we are doing atomic operations, that variable should be created using **atomic_t** or **atomic64_t**. So we have separate special functions for reading, writing, and arithmetic operations, and those are explained below.

The declarations are needed to use the atomic integer operations are in **<asm/atomic.h>**. Some architectures provide additional methods that are unique to that architecture, but all architectures provide at least a minimum set of operations that are used throughout the kernel. When you write kernel code, you can ensure that these operations are correctly implemented on all architectures.

Creating atomic variables

```
1 atomic_t etx_global_variable; /* define etx_global_variable */
2
3 or
4
5 atomic_t etx_global_variable = ATOMIC_INIT(0); /* define etx_global_variable and init
```

Reading atomic variables

atomic_read

This function atomically reads the value of the given atomic variable.

```
int atomic_read(atomic_t *v);
```

where,

v - pointer of type atomic_t

Return: It returns the integer value.

Other operations on atomic variables

atomic_set

This function atomically sets the value to the atomic variable.


```
void atomic_set(atomic_t *v, int i);
```

where,

v – the pointer of type atomic_t

i – the value to be set to v

atomic_add

This function atomically adds the value to the atomic variable.

```
void atomic_add(int i, atomic_t *v);
```

where,

i – the value to be added to v

v – the pointer of type atomic_t

atomic_sub

This function atomically subtracts the value from the atomic variable.

```
void atomic_sub(int i, atomic_t *v);
```

where,

i – the value to be subtracted from v

v – the pointer of type atomic_t

atomic_inc

This function atomically increments the value of the atomic variable by

1.

```
void atomic_inc (atomic_t *v);
```

where,

v - the pointer of type atomic_t

atomic_dec

This function atomically decrements the value of the atomic variable by 1.

```
void atomic_dec (atomic_t *v);
```

where,

v - the pointer of type atomic_t

atomic_sub_and_test

This function atomically subtract the value from the atomic variable and test the result is zero or not.

```
void atomic_sub_and_test(int i, atomic_t *v);
```

where,

i - the value to be subtracted from v

v - the pointer of type atomic_t

Return: It returns true if the result is zero, or false for all other cases.

atomic_dec_and_test

This function atomically decrements the value of the atomic variable by 1 and test the result is zero or not.

1

```
void atomic_dec_and_test(atomic_t *v);
```

where,

v - the pointer of type atomic_t

Return: It returns true if the result is zero, or false for all other cases.

atomic_inc_and_test

This function atomically increments the value of the atomic variable by 1 and test the result is zero or not.

```
void atomic_inc_and_test(atomic_t *v);
```

where,

v - the pointer of type atomic_t

Return: It returns true if the result is zero, or false for all other cases.

atomic_add_negative

This function atomically adds the value to the atomic variable and test the result is negative or not.

```
void atomic_add_negative(int i, atomic_t *v);
```

where,

i - the value to be added to v

v - the pointer of type atomic_t

Return: It returns true if the result is negative, or false for all other cases.

atomic_add_return

This function atomically adds the value to the atomic variable and returns the value.

```
void atomic_add_return(int i, atomic_t *v);
```

1

where,

i - the value to be added to v

v - the pointer of type atomic_t

Return : It returns true if the result the value (i + v).

Like this, other functions also there. Those are,

Function	Description
<code>int atomic_sub_return(int i, atomic_t *v)</code>	Atomically subtract i from v and return the result
<code>int atomic_inc_return(int i, atomic_t *v)</code>	Atomically increments v by one and return the result
<code>int atomic_dec_return(int i, atomic_t *v)</code>	Atomically decrements v by one and return the result

atomic_add_unless

This function atomically adds the value to the atomic variable unless the number is a given value.

```
atomic_add_unless (atomic_t *v, int a, int u);
```

where,

v – the pointer of type `atomic_t`

a – the amount to add to **v**...

u – ...unless **v** is equal to **u**.

Return: It returns non-zero if **v** was not **u**, and zero otherwise.

There is a 64-bit version also available. Unlike `atomic_t`, that will be operates on 64 bits. This 64-bit version also have similar function like above, the only change is we have to use 64.

Example

```
1 atomic64_t etx_global_variable = ATOMIC64_INIT(0);
2 long atomic64_read(atomic64_t *v);
3 void atomic64_set(atomic64_t *v, int i);
4 void atomic64_add(int i, atomic64_t *v);
5 void atomic64_sub(int i, atomic64_t *v);
6 void atomic64_inc(atomic64_t *v);
7 void atomic64_dec(atomic64_t *v);
8 int atomic64_sub_and_test(int i, atomic64_t *v);
```

```

9  int atomic64_add_negative(int i, atomic64_t *v);
10 long atomic64_add_return(int i, atomic64_t *v);
11 long atomic64_sub_return(int i, atomic64_t *v);
12 long atomic64_inc_return(int i, atomic64_t *v);
13 long atomic64_dec_return(int i, atomic64_t *v);
14 int atomic64_dec_and_test(atomic64_t *v);
15 int atomic64_inc_and_test(atomic64_t *v);

```

But all the operations are the same as **atomic_t**.

Atomic Bitwise Operations

Atomic_t is good when we are working on integer arithmetic. But when it comes to bitwise atomic operation, it doesn't work well. So kernel offers separate functions to achieve that. Atomic bit operations are very fast. These functions are architecture dependent and are declared in **<asm/bitops.h>**.

These bitwise functions operate on a generic pointer. So **atomic_t** / **atomic64_t** is not required. So we can work with a pointer to whatever data we want.

The below functions are available for atomic bit operations.

Function	Description
void set_bit(int nr, void *addr)	Atomically set the nr -th bit starting from addr
void clear_bit(int nr, void *addr)	Atomically clear the nr -th bit starting from addr
void change_bit(int nr, void *addr)	Atomically flip the value of the nr -th bit starting from addr
int test_and_set_bit(int nr, void *addr)	Atomically set the nr -th bit starting from addr and return the previous value
int test_and_clear_bit(int nr, void *addr)	Atomically clear the nr -th bit starting from addr and return the previous value

int test_and_change_bit(int nr, void *addr)	Atomically flip the nr -th bit starting from addr and return the previous value
int test_bit(int nr, void *addr)	Atomically return the value of the nr -th bit starting from addr
int find_first_zero_bit(unsigned long *addr, unsigned int size)	Atomically returns the bit-number of the first zero bit, not the number of the byte containing a bit
int find_first_bit(unsigned long *addr, unsigned int size)	Atomically returns the bit-number of the first set bit, not the number of the byte containing a bit

And also non-atomic bit operations also available. What is the use of that when we have atomic bit operations? When we have code which is already locked by **mutex/spinlock** then we can go for this non-atomic version. This might be faster in that case. The below functions are available for non-atomic bit operations.

Function	Description
void _set_bit(int nr, void *addr)	Non-atomically set the nr -th bit starting from addr
void _clear_bit(int nr, void *addr)	Non-atomically clear the nr -th bit starting from addr
void _change_bit(int nr, void *addr)	Non-atomically flip the value of the nr -th bit starting from addr
int _test_and_set_bit(int nr, void *addr)	Non-atomically set the nr -th bit starting from addr and return the previous value
int _test_and_clear_bit(int nr, void *addr)	Non-atomically clear the nr -th bit starting from addr and return the previous value
int _test_and_change_bit(int nr, void *addr)	Non-atomically flip the nr -th bit starting from addr and return the

	previous value
int _test_bit(int nr, void *addr)	Non-atomically return the value of the nr -th bit starting from addr

Example Programming

In this program, we have two threads called **thread_function1** and **thread_function2**. Both will be accessing the atomic variables.

Driver Source Code

```

1  #include <linux/kernel.h>
2  #include <linux/init.h>
3  #include <linux/module.h>
4  #include <linux/kdev_t.h>
5  #include <linux/fs.h>
6  #include <linux/cdev.h>
7  #include <linux/device.h>
8  #include <linux/slab.h>           //kmalloc()
9  #include <linux/uaccess.h>       //copy_to/from_user()
10 #include <linux/kthread.h>       //kernel threads
11 #include <linux/sched.h>         //task_struct
12 #include <linux/delay.h>
13
14 atomic_t etx_global_variable = ATOMIC_INIT(0);    //Atomic integer variable
15 unsigned int etc_bit_check = 0;
16
17 dev_t dev = 0;
18 static struct class *dev_class;
19 static struct cdev etx_cdev;
20
21 static int __init etx_driver_init(void);
22 static void __exit etx_driver_exit(void);
23
24 static struct task_struct *etx_thread1;
25 static struct task_struct *etx_thread2;
26
27 /***** Driver Fuctions *****/
28 static int etx_open(struct inode *inode, struct file *file);
29 static int etx_release(struct inode *inode, struct file *file);
30 static ssize_t etx_read(struct file *filp,
31                        char __user *buf, size_t len, loff_t * off);
32 static ssize_t etx_write(struct file *filp,
33                        const char *buf, size_t len, loff_t * off);
34 /*****/
35
36 int thread_function1(void *pv);
37 int thread_function2(void *pv);
38
39 int thread_function1(void *pv)
40 {
41     unsigned int prev_value = 0;
42
43     while(!kthread_should_stop()) {
44         atomic_inc(&etx_global_variable);

```

```

45     prev_value = test_and_change_bit(1, (void*)&etc_bit_check);
46     printk(KERN_INFO "Function1 [value : %u] [bit:%u]\n", atomic_read(&etx_glob
47     msleep(1000);
48 }
49 return 0;
50 }
51
52 int thread_function2(void *pv)
53 {
54     unsigned int prev_value = 0;
55     while(!kthread_should_stop()) {
56         atomic_inc(&etx_global_variable);
57         prev_value = test_and_change_bit(1,(void*) &etc_bit_check);
58         printk(KERN_INFO "Function2 [value : %u] [bit:%u]\n", atomic_read(&etx_glob
59         msleep(1000);
60     }
61     return 0;
62 }
63
64 static struct file_operations fops =
65 {
66     .owner          = THIS_MODULE,
67     .read           = etx_read,
68     .write          = etx_write,
69     .open           = etx_open,
70     .release        = etx_release,
71 };
72
73 static int etx_open(struct inode *inode, struct file *file)
74 {
75     printk(KERN_INFO "Device File Opened...!!!\n");
76     return 0;
77 }
78
79 static int etx_release(struct inode *inode, struct file *file)
80 {
81     printk(KERN_INFO "Device File Closed...!!!\n");
82     return 0;
83 }
84
85 static ssize_t etx_read(struct file *filp,
86                         char __user *buf, size_t len, loff_t *off)
87 {
88     printk(KERN_INFO "Read function\n");
89     return 0;
90 }
91
92 static ssize_t etx_write(struct file *filp,
93                          const char __user *buf, size_t len, loff_t *off)
94 {
95     printk(KERN_INFO "Write Function\n");
96     return len;
97 }
98
99 static int __init etx_driver_init(void)
100 {
101     /*Allocating Major number*/
102     if((alloc_chrdev_region(&dev, 0, 1, "etx_Dev")) <0){
103         printk(KERN_INFO "Cannot allocate major number\n");
104         return -1;
105     }
106     printk(KERN_INFO "Major = %d Minor = %d \n",MAJOR(dev), MINOR(dev));
107

```



```
108      /*Creating cdev structure*/
109      cdev_init(&etx_cdev,&fops);
110
111      /*Adding character device to the system*/
112      if((cdev_add(&etx_cdev,dev,1)) < 0){
113          printk(KERN_INFO "Cannot add the device to the system\n");
114          goto r_class;
115      }
116
117      /*Creating struct class*/
118      if((dev_class = class_create(THIS_MODULE,"etx_class")) == NULL){
119          printk(KERN_INFO "Cannot create the struct class\n");
120          goto r_class;
121      }
122
123      /*Creating device*/
124      if((device_create(dev_class,NULL,dev,NULL,"etx_device")) == NULL){
125          printk(KERN_INFO "Cannot create the Device \n");
126          goto r_device;
127      }
128
129
130      /* Creating Thread 1 */
131      etx_thread1 = kthread_run(thread_function1,NULL,"eTx Thread1");
132      if(etx_thread1) {
133          printk(KERN_ERR "Kthread1 Created Successfully...\n");
134      } else {
135          printk(KERN_ERR "Cannot create kthread1\n");
136          goto r_device;
137      }
138
139      /* Creating Thread 2 */
140      etx_thread2 = kthread_run(thread_function2,NULL,"eTx Thread2");
141      if(etx_thread2) {
142          printk(KERN_ERR "Kthread2 Created Successfully...\n");
143      } else {
144          printk(KERN_ERR "Cannot create kthread2\n");
145          goto r_device;
146      }
147
148      printk(KERN_INFO "Device Driver Insert...Done!!!\n");
149      return 0;
150
151
152  r_device:
153      class_destroy(dev_class);
154  r_class:
155      unregister_chrdev_region(dev,1);
156      cdev_del(&etx_cdev);
157      return -1;
158  }
159
160  void __exit etx_driver_exit(void)
161  {
162      kthread_stop(etx_thread1);
163      kthread_stop(etx_thread2);
164      device_destroy(dev_class,dev);
165      class_destroy(dev_class);
166      cdev_del(&etx_cdev);
167      unregister_chrdev_region(dev, 1);
168      printk(KERN_INFO "Device Driver Remove...Done!!!\n");
169  }
170
```

```
171 module_init(etx_driver_init);
172 module_exit(etx_driver_exit);
173
174 MODULE_LICENSE("GPL");
175 MODULE_AUTHOR("EmbeTronicX <embetronicx@gmail.com>");
176 MODULE_DESCRIPTION("A simple device driver - Atomic Variables");
177 MODULE_VERSION("1.27");
```

MakeFile

```
1 obj-m += driver.o
2
3 KDIR = /lib/modules/$(shell uname -r)/build
4
5 all:
6     make -C $(KDIR) M=$(shell pwd) modules
7
8 clean:
9     make -C $(KDIR) M=$(shell pwd) clean
```

In our [next tutorial](#), we can discuss Seqlock.

0

Article Rating



Share this:

Share 7

Post

Tweet

SHARE

Print

WhatsApp

Share

1

Email

Telegram

More

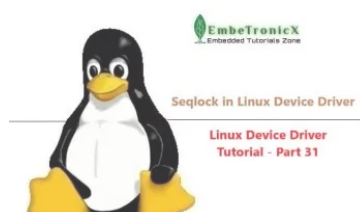
Like this:

Loading...

Related



Linux Device Driver



Linux Device Driver



Linux Device Driver

[Tutorial Part 23 – Spinlock
in Linux Kernel Part 1](#)
In "Device Drivers"

[Tutorial Part 31 – Seqlock
in Linux Kernel](#)
In "Device Drivers"

[Tutorial Part 24 – Read
Write Spinlock in Linux
Kernel \(Spinlock Part 2\)](#)
In "Device Drivers"

1



☒ Subscribe ▼

Connect with | [Login](#)

